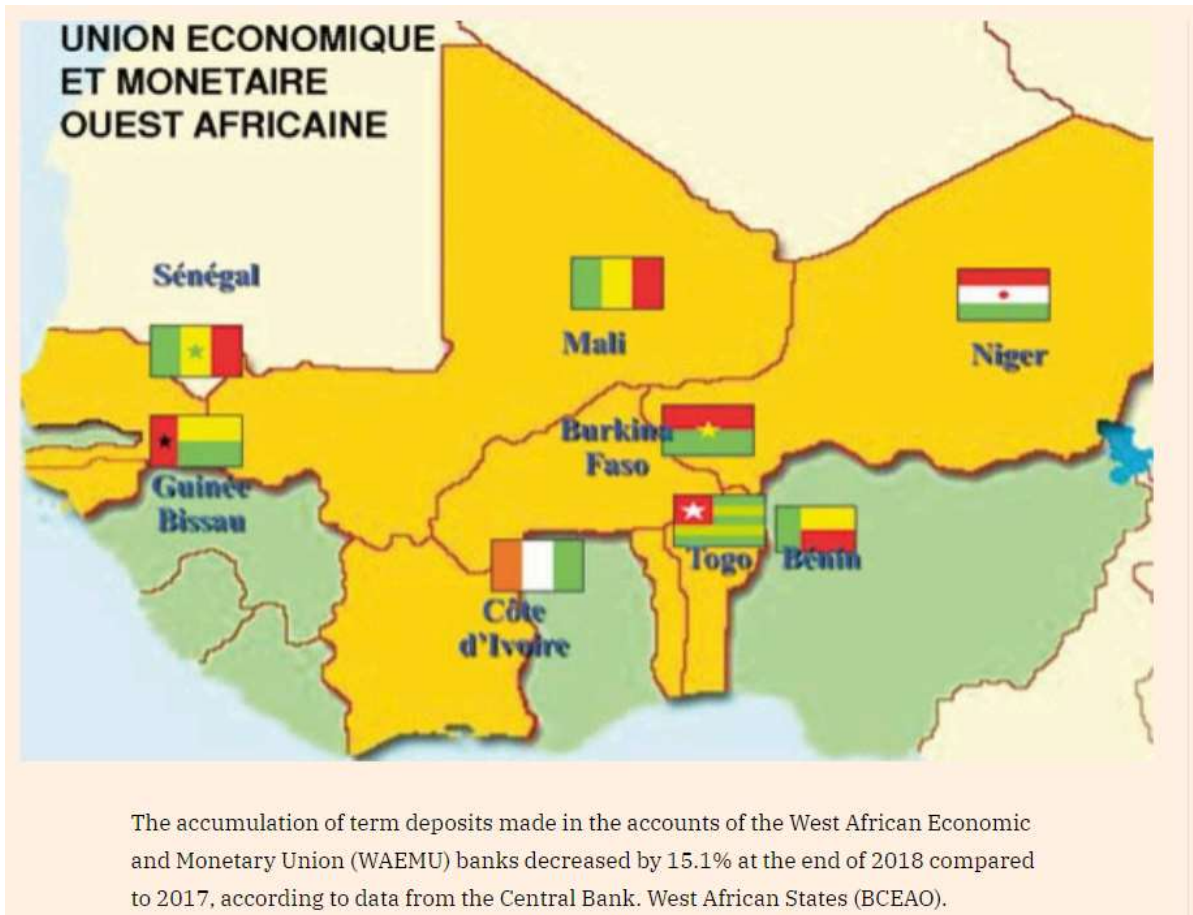


```
In [1]: 1 from IPython.display import Image, display
2
3 # Specify the path to your image file
4 image_path = r'C:\Users\jangi\Downloads\Screenshot 2024-05-15 200823.png'
5
6 # Display the image
7 display(Image(filename=image_path))
```



Insight of this dataset:

1. Predicting Financial Health or Bankruptcy (e.g., using Zscore):

- Logistic Regression: Suitable for binary classification (e.g., predicting whether a bank will go bankrupt).
- Random Forest: Effective for classification problems and can handle non-linear relationships.
- Gradient Boosting (e.g., XGBoost): Provides high performance for classification tasks and handles various complexities in the data.
- Neural Networks: Useful for complex patterns but requires more data and computational power.

2. Assessing Profitability and Sustainability (PS):

- Linear Regression: Simple and interpretable, good for predicting continuous outcomes.
- Random Forest Regression: Handles non-linear relationships and provides variable importance measures.
- Gradient Boosting Regression: Offers high accuracy and can model complex patterns.
- Neural Networks: Suitable for capturing complex dependencies in large datasets.

3. Solvency and Financial Stability (SFS):

- Linear Regression or Logistic Regression: Depending on whether SFS is a continuous or binary variable.
- Support Vector Machines (SVM): Effective for both regression and classification, especially in high-dimensional spaces.
- Random Forest/Gradient Boosting: Versatile and effective for both regression and classification.

4. Liquidity Risk (RL):

- Logistic Regression: If predicting the probability of high liquidity risk.
- Random Forest/Gradient Boosting: Useful for classification tasks and handling various feature types.
- Time Series Analysis (e.g., ARIMA, LSTM): If liquidity risk is being predicted over time.

5. Capital Adequacy (CC):

- Linear Regression: For predicting continuous capital adequacy ratios.
- Random Forest/Gradient Boosting Regression: Effective for non-linear relationships.
- Support Vector Regression (SVR): Good for continuous predictions with non-linear relationships.
- Algorithm Selection Based on Data Characteristics
- Logistic Regression: Best for binary outcomes and interpretability.
- Random Forest/Gradient Boosting: Good for handling non-linear relationships and feature interactions.
- Neural Networks: Suitable for large datasets with complex patterns but require more computational resources.
- Time Series Models: Necessary if predicting variables over time.
- Example Workflow for Predicting Financial Health (using Zscore)

Step-by-Step Process

- Import Libraries

```
In [2]: 1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
5 from sklearn.metrics import classification_report, confusion_matrix
6 from sklearn.metrics import mean_squared_error, r2_score
7 from sklearn.preprocessing import StandardScaler, OneHotEncoder
8 from sklearn.compose import ColumnTransformer
9 from sklearn.pipeline import Pipeline
10 import matplotlib.pyplot as plt
11 import seaborn as sns
12 sns.set()
13 %matplotlib inline
14 import warnings
15 warnings.filterwarnings("ignore")
16
```

```
In [3]: 1 df = pd.read_csv(r"WAEMU_Banking.csv")
```

```
In [4]: 1 df = df.loc[:, ~df.columns.str.contains('^Unnamed')]
```

```
In [5]: 1 df.head()
```

```
Out[5]:
```

	Countries_Num	id	Countries	Banks	Year	RIR	SFS	INF	ERA	
0	1	1	2	13	2013	3.836593	26.861971	0.428889	3.196428	12.076
1	1	1	2	13	2014	5.599992	29.965430	-0.548758	3.045024	8.884
2	1	1	2	13	2015	4.266334	30.984761	0.218786	2.394557	8.583
3	1	1	2	13	2016	4.580100	29.832095	-0.794050	3.712403	5.720
4	1	1	2	13	2017	7.329021	28.630991	1.769412	3.833422	6.256

Data Preprocessing:

- Handle missing values.
- Normalize/standardize numerical features.
- Encode categorical variables.

Feature Selection/Engineering:

- Select relevant features that contribute to the target variable.
- Create new features if necessary (e.g., ratios, interaction terms).

Model Selection:

- Choose an appropriate algorithm (e.g., Random Forest, Gradient Boosting).

Model Training and Evaluation:

- Split data into training and testing sets.
- Train the model on the training set.
- Evaluate performance using metrics like accuracy, precision, recall, F1 score (for classification) or RMSE, MAE (for regression).

Hyperparameter Tuning:

- Use techniques like Grid Search or Random Search to find the best hyperparameters.

Model Validation:

- Validate the model using cross-validation to ensure it generalizes well to unseen data.

In [6]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 742 entries, 0 to 741
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Countries_Num    742 non-null    int64
1   id               742 non-null    int64
2   Countries        742 non-null    int64
3   Banks            742 non-null    int64
4   Year             742 non-null    int64
5   RIR              742 non-null    float64
6   SFS              742 non-null    float64
7   INF              742 non-null    float64
8   ERA              742 non-null    float64
9   INL              742 non-null    float64
10  Zscore           742 non-null    float64
11  DEBT             742 non-null    float64
12  SIZE             742 non-null    float64
13  CC               742 non-null    float64
14  GE               742 non-null    float64
15  PS               742 non-null    float64
16  RQ               742 non-null    float64
17  RL               742 non-null    float64
18  VA               742 non-null    float64
dtypes: float64(14), int64(5)
memory usage: 110.3 KB
```

```
In [7]: 1 df.shape
```

```
Out[7]: (742, 19)
```

```
In [8]: 1 df.isnull().sum()
```

```
Out[8]: Countries_Num    0  
id                      0  
Countries              0  
Banks                  0  
Year                   0  
RIR                     0  
SFS                     0  
INF                     0  
ERA                     0  
INL                     0  
Zscore                  0  
DEBT                    0  
SIZE                    0  
CC                      0  
GE                      0  
PS                      0  
RQ                      0  
RL                      0  
VA                      0  
dtype: int64
```

In [9]: 1 df.describe().T.style.background_gradient(cmap = 'Reds')

Out[9]:

	count	mean	std	min	25%	50%	
Countries_Num	742.000000	4.575472	2.316191	1.000000	3.000000	5.000000	
id	742.000000	53.500000	30.618842	1.000000	27.000000	53.500000	8
Countries	742.000000	3.650943	2.380928	0.000000	2.000000	3.000000	
Banks	742.000000	48.716981	27.703134	0.000000	26.000000	47.500000	7
Year	742.000000	2016.000000	2.001349	2013.000000	2014.000000	2016.000000	201
RIR	742.000000	3.711618	4.114070	-23.137938	3.367054	4.266334	
SFS	742.000000	31.965249	8.239127	15.829639	26.927042	29.918364	3
INF	742.000000	0.524620	1.269270	-3.233389	-0.258090	0.685881	
ERA	742.000000	9.766718	18.965490	-179.747455	4.166364	7.460494	1
INL	742.000000	11.652364	10.885343	0.000000	4.656050	8.895584	1
Zscore	742.000000	2.967964	5.123174	-47.777093	1.482417	2.149078	
DEBT	742.000000	38.780738	11.664036	18.503746	30.666445	36.494147	4
SIZE	742.000000	12.077644	1.114687	8.677440	11.359607	12.156868	1
CC	742.000000	37.292780	12.589451	13.461540	27.403850	33.173080	4
GE	742.000000	26.747784	10.261396	8.530806	18.009480	26.442310	3
PS	742.000000	23.891205	14.741995	3.773585	12.857140	17.061610	3
RQ	742.000000	35.432019	9.474103	13.942310	28.365390	34.134620	4
RL	742.000000	32.653598	10.138651	6.250000	27.403850	30.769230	3
VA	742.000000	41.074591	11.071831	18.309860	33.004920	38.423650	5



In [10]: 1 df.duplicated()

Out[10]:

```


0      False
1      False
2      False
3      False
4      False
...
737    False
738    False
739    False
740    False
741    False
Length: 742, dtype: bool

```



```
In [11]: 1 df.nunique().sort_values(ascending=True)
```

```
Out[11]: Year          7
Countries_Num        8
Countries            8
RQ                  42
CC                  42
RL                  42
GE                  47
VA                  47
PS                  50
SFS                 56
DEBT                56
RIR                 56
INF                 56
Banks               98
id                 106
ERA                 711
INL                 712
SIZE                714
Zscore              718
dtype: int64
```

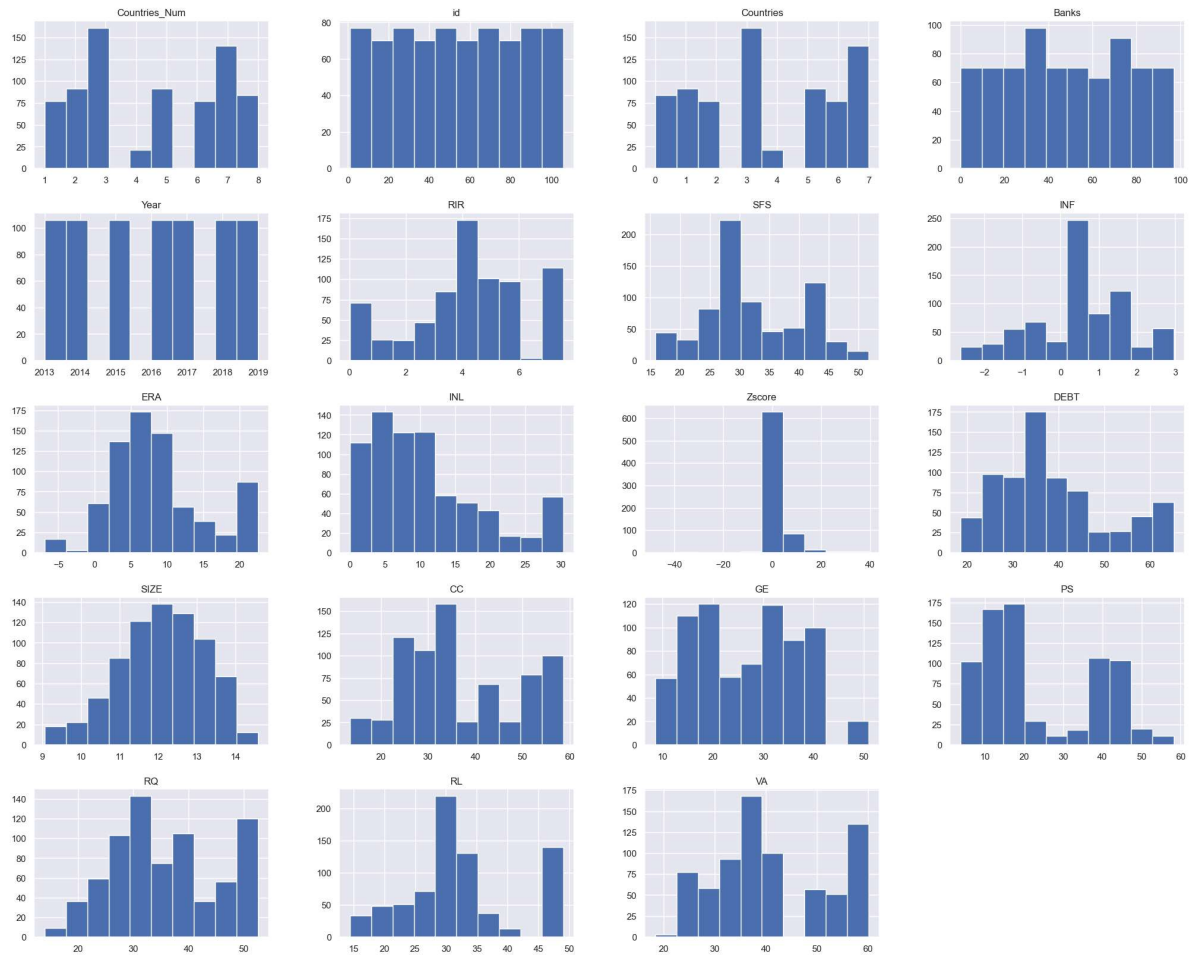
```
In [12]: 1 def boxplots(col):
2         plt.figure(figsize=(5,4))
3         sns.boxplot(df,x=col,palette='husl')
4         plt.show()
5
6         # Assuming 'df' is your DataFrame
7         for i in list(df.select_dtypes(exclude=['object']).columns)[0:]:
8             boxplots(i)
```



```
In [13]: 1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5
6 # Load your own dataset
7 # Replace 'your_dataset.csv' with the actual file path or URL of your data
8 df = df
9
10 # Function to handle outliers using clip
11 def handle_outliers(dataframe, col):
12     q3 = dataframe[col].quantile(0.75)
13     q1 = dataframe[col].quantile(0.25)
14     IQR = q3 - q1
15     Lower = q1 - 1.5 * IQR
16     Upper = q3 + 1.5 * IQR
17
18     # Clip the values to be within the Lower and Upper bounds
19     dataframe[col] = dataframe[col].clip(lower=Lower, upper=Upper)
20
21 # Specify the columns to handle outliers
22 columns_to_handle_outliers = [
23     'RIR',
24     'INF',
25     'ERA',
26     'INL',
27     'DEBT',
28     'SIZE',
29     'RL',
30 ]
31
32 # Plot boxplots before handling outliers
33 plt.figure(figsize=(30, 6))
34 plt.subplot(1, 2, 1)
35 sns.boxplot(data=df[columns_to_handle_outliers].dropna())
36 plt.title('Boxplots Before Handling Outliers')
37
38 # Handle outliers for each specified column
39 for col in columns_to_handle_outliers:
40     handle_outliers(df, col)
41
42 # Plot boxplots after handling outliers
43 plt.subplot(1, 2, 2)
44 sns.boxplot(data=df[columns_to_handle_outliers].dropna())
45 plt.title('Boxplots After Handling Outliers')
46
47 plt.show()
48
```

```
In [14]: 1 df.hist(figsize = (25,20))
          2 plt.show()
```



```
In [15]: 1 # Exploratory Data Analysis (EDA)
2
3 # Setting aesthetics for plots
4 sns.set_style('whitegrid')
5
6 # A. Descriptive Statistics
7 print("Descriptive Statistics:")
8 print(df.describe())
```

Descriptive Statistics:

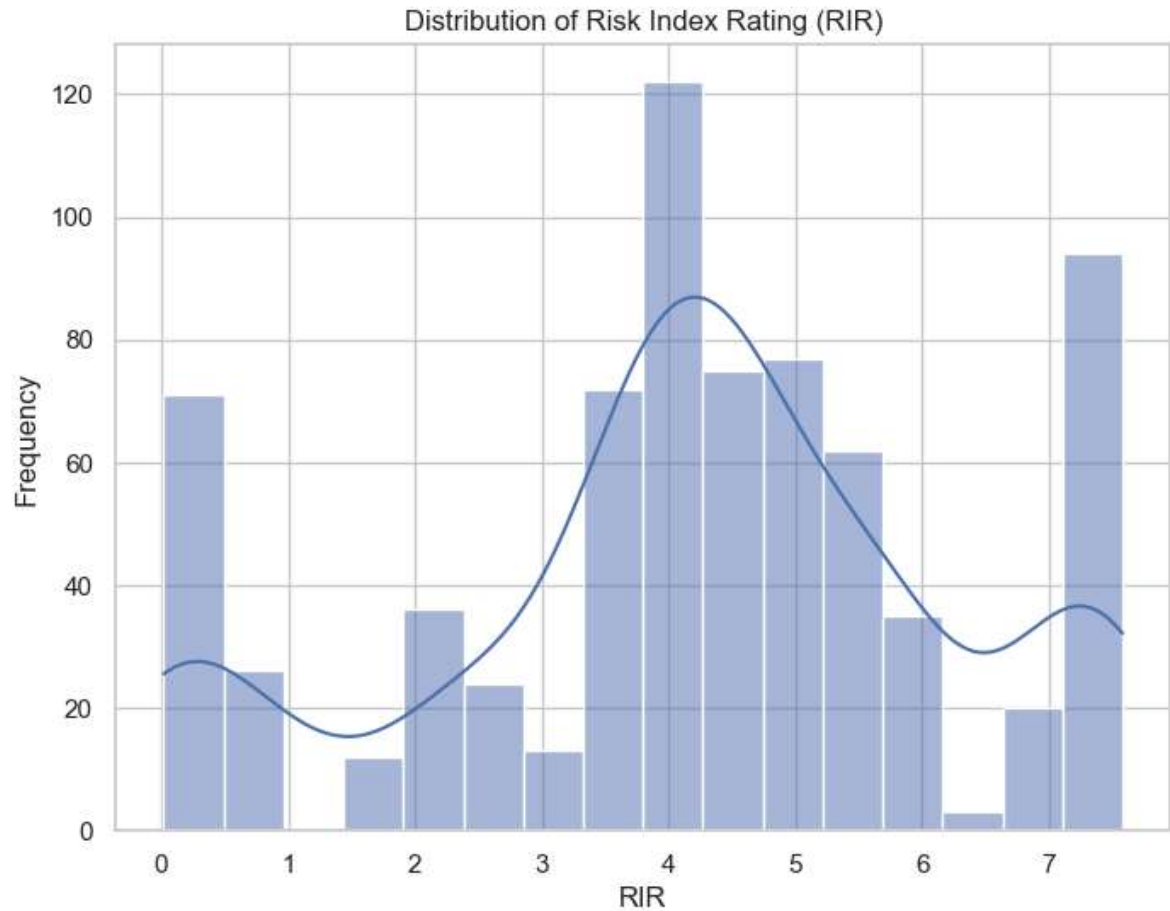
	Countries_Num	id	Countries	Banks	Year	\
count	742.000000	742.000000	742.000000	742.000000	742.000000	
mean	4.575472	53.500000	3.650943	48.716981	2016.000000	
std	2.316191	30.618842	2.380928	27.703134	2.001349	
min	1.000000	1.000000	0.000000	0.000000	2013.000000	
25%	3.000000	27.000000	2.000000	26.000000	2014.000000	
50%	5.000000	53.500000	3.000000	47.500000	2016.000000	
75%	7.000000	80.000000	6.000000	74.000000	2018.000000	
max	8.000000	106.000000	7.000000	97.000000	2019.000000	

	RIR	SFS	INF	ERA	INL	Zscore
\						
count	742.000000	742.000000	742.000000	742.000000	742.000000	742.000000
mean	4.180301	31.965249	0.535323	8.683170	10.825335	2.967964
std	2.065988	8.239127	1.239726	6.636001	8.166584	5.123174
min	0.017647	15.829639	-2.622454	-6.865332	0.000000	-47.777093
25%	3.367054	26.927042	-0.258090	4.166364	4.656050	1.482417
50%	4.266334	29.918364	0.685881	7.460494	8.895584	2.149078
75%	5.599992	37.793417	1.318153	11.520829	14.951115	3.468952
max	7.578020	51.682209	2.967604	22.552525	30.393713	39.380715

	DEBT	SIZE	CC	GE	PS	RQ
\						
count	742.000000	742.000000	742.000000	742.000000	742.000000	742.000000
mean	38.778321	12.078829	37.292780	26.747784	23.891205	35.432019
std	11.658474	1.111316	12.589451	10.261396	14.741995	9.474103
min	18.503746	9.061235	13.461540	8.530806	3.773585	13.942310
25%	30.666445	11.359607	27.403850	18.009480	12.857140	28.365390
50%	36.494147	12.156868	33.173080	26.442310	17.061610	34.134620
75%	44.508726	12.891855	49.038460	35.096150	38.571430	43.269230
max	65.272147	14.582210	58.653850	50.961540	58.293840	52.606640

	RL	VA
count	742.000000	742.000000
mean	32.523353	41.074591
std	9.209518	11.071831
min	14.423090	18.309860
25%	27.403850	33.004920
50%	30.769230	38.423650
75%	36.057690	53.490507
max	49.038450	60.098520

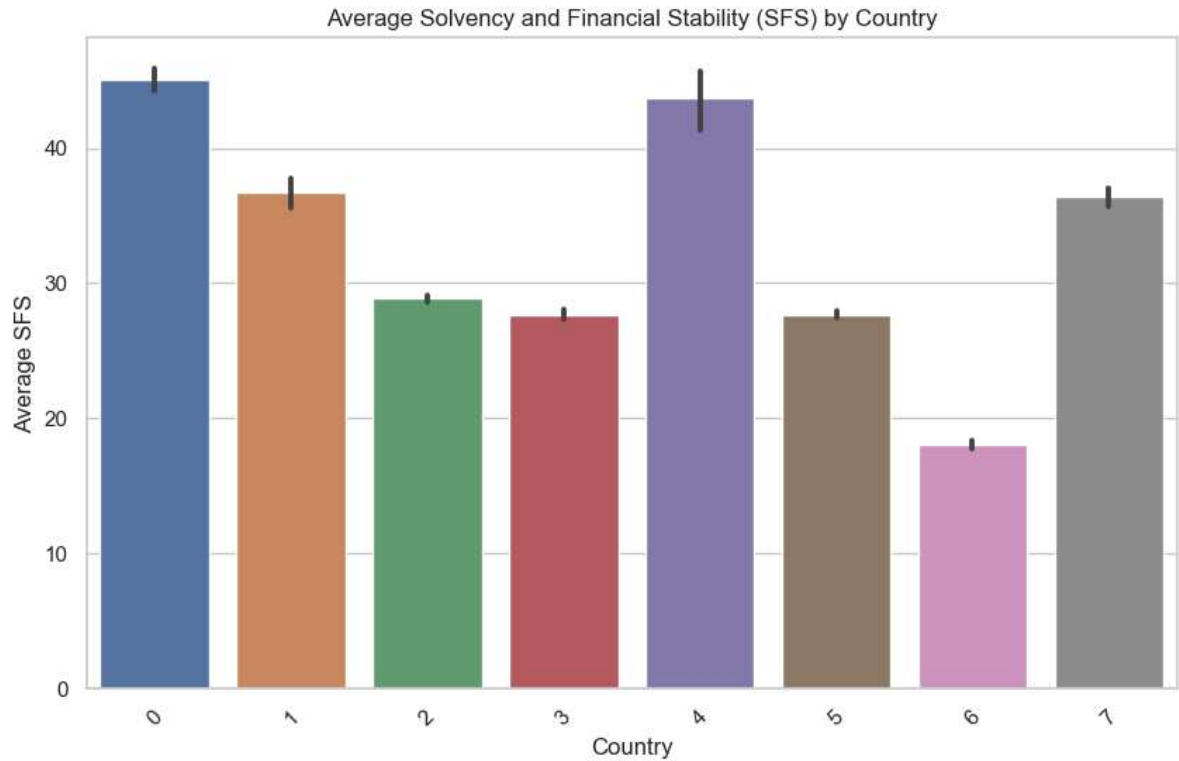
```
In [16]: 1 # B. Distribution of Key Variables
2 # Example: Distribution of Risk Index Rating (RIR)
3 plt.figure(figsize=(8, 6))
4 sns.histplot(df['RIR'], kde=True)
5 plt.title('Distribution of Risk Index Rating (RIR)')
6 plt.xlabel('RIR')
7 plt.ylabel('Frequency')
8 plt.show()
```



```

In [17]: 1 # C. Comparison of Metrics Across Different Countries
2 # Example: Average Solvency and Financial Stability (SFS) by Country
3 plt.figure(figsize=(10, 6))
4 sns.barplot(x='Countries', y='SFS', data=df)
5 plt.title('Average Solvency and Financial Stability (SFS) by Country')
6 plt.xlabel('Country')
7 plt.ylabel('Average SFS')
8 plt.xticks(rotation=45)
9 plt.show()


```




```

In [18]: 1 plt.figure(figsize=(14, 6))
2 sns.kdeplot(data=df, x='VA',
3             hue='Countries',          # to differentiate between the different countries
4             fill=True,                # to fill the area under the curve
5             alpha=0.5, linewidth=1.5 # to make the plot more visually appealing
6             )
7
8
9 # Add a title and labels to the plot using Matplotlib
10 plt.title("Density Plot of VA by Countries")
11 plt.xlabel("VA")
12 plt.ylabel("Density")
13
14 # display the plot
15 plt.show()

```





```
In [19]: 1 sns.jointplot(data=df,          # Plot bivariate distribution
2             x='RL', y='VA',
3             kind='reg')          # kind= reg, kde, hex
4
5 # display the plot
6 plt.show()
```



```
In [20]: 1 # Let's see what happens if we change the kind to kde
2 sns.jointplot(x='RL', y='VA', data=df, kind='kde')
3
4 # Let's see what happens if we change the kind to hex
5 sns.jointplot(x='RL', y='VA', data=df, kind='hex')
```

Out[20]: <seaborn.axisgrid.JointGrid at 0x155c762bbb0>





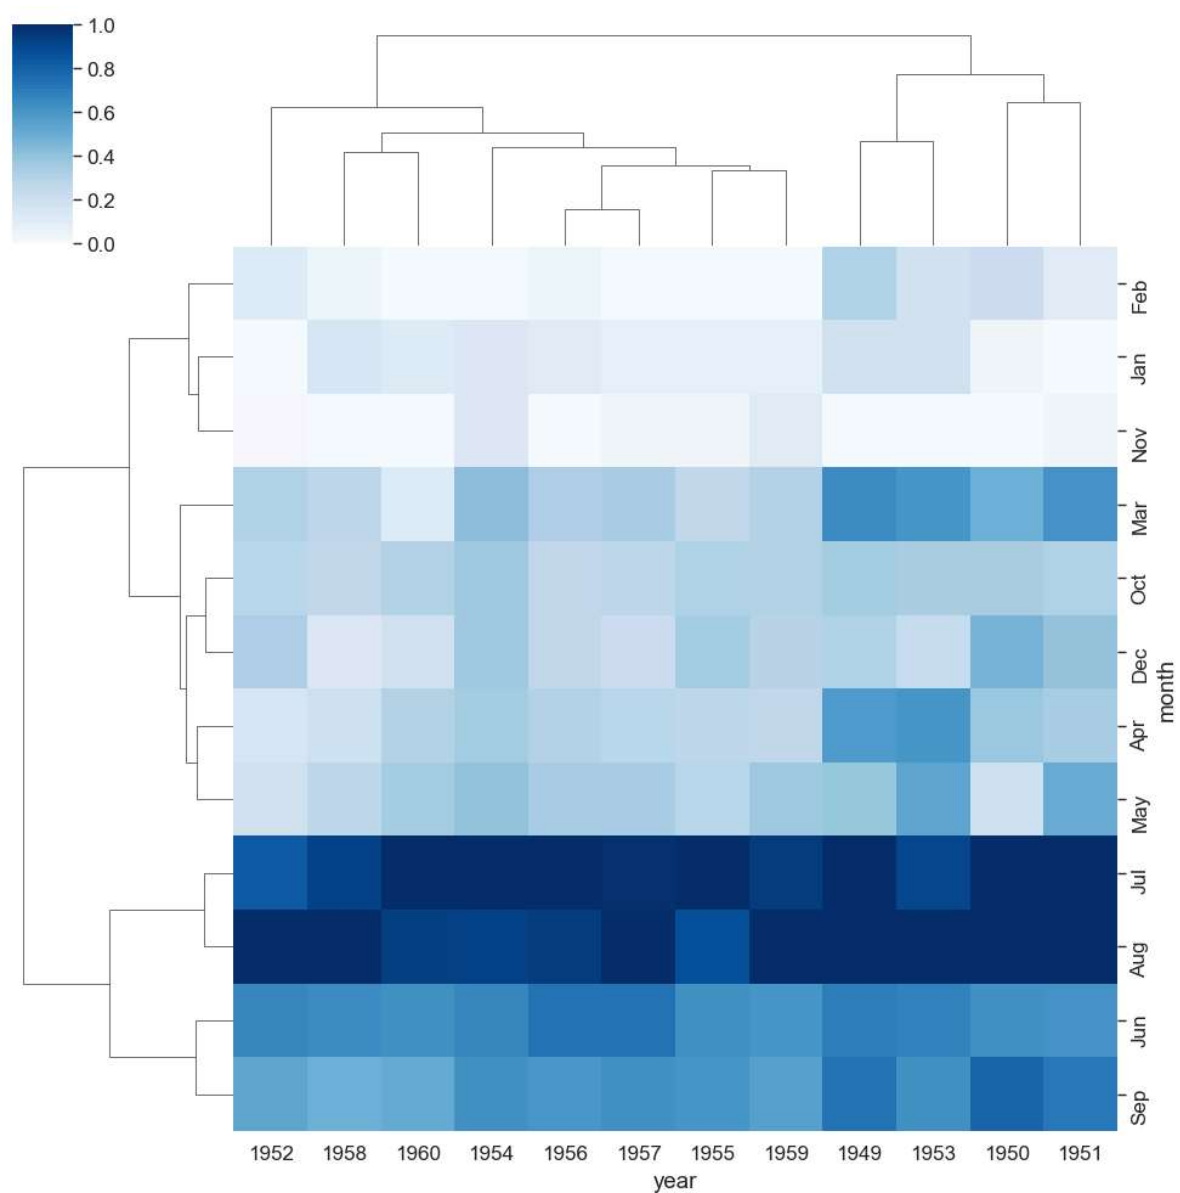

```

In [21]: 1 plt.figure(figsize=(4,6))
          2 sns.set_context('paper', font_scale=1.4)
          3
          4 # let's use flights dataset
          5 flights = sns.load_dataset("flights")
          6 flights = flights.pivot_table(index='month', columns='year', values='passenger')
          7
          8 sns.clustermap(flights, cmap="Blues", standard_scale=1)

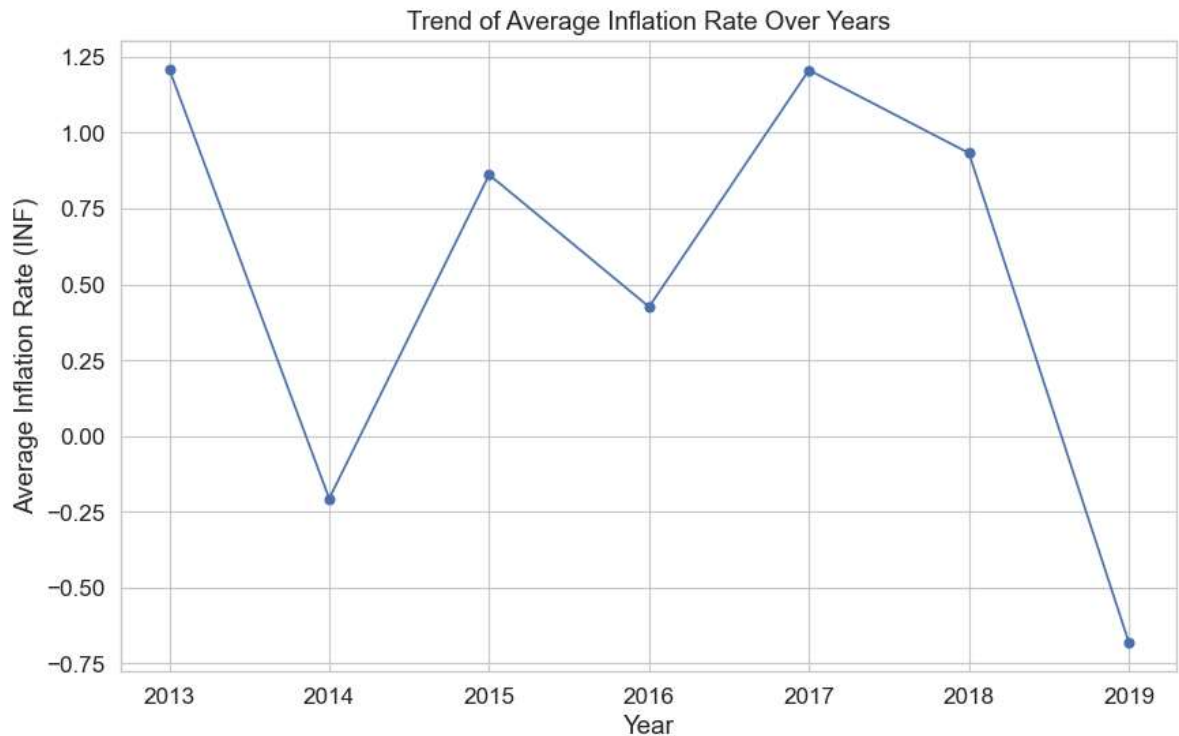
```

Out[21]: <seaborn.matrix.ClusterGrid at 0x155c736b040>

<Figure size 400x600 with 0 Axes>



```
In [22]: 1 # D. Time-Series Analysis for Yearly Data
2 # Example: Trend of Average Inflation Rate (INF) Over Years
3 plt.figure(figsize=(10, 6))
4 df.groupby('Year')['INF'].mean().plot(kind='line', marker='o')
5 plt.title('Trend of Average Inflation Rate Over Years')
6 plt.xlabel('Year')
7 plt.ylabel('Average Inflation Rate (INF)')
8 plt.show()
```



```
In [23]: 1 df['Countries'] = df['Countries'].astype('category')
2 df['Countries'] = df['Countries'].cat.codes
3 df.head()
```

```
Out[23]:
```

	Countries_Num	id	Countries	Banks	Year	RIR	SFS	INF	ERA
0	1	1	2	13	2013	3.836593	26.861971	0.428889	3.196428
1	1	1	2	13	2014	5.599992	29.965430	-0.548758	3.045024
2	1	1	2	13	2015	4.266334	30.984761	0.218786	2.394557
3	1	1	2	13	2016	4.580100	29.832095	-0.794050	3.712403
4	1	1	2	13	2017	7.329021	28.630991	1.769412	3.833422

```
In [24]: 1 df['Banks'] = df['Banks'].astype('category')
2 df['Banks'] = df['Banks'].cat.codes
3 df.head()
```

```
Out[24]:
```

	Countries_Num	id	Countries	Banks	Year	RIR	SFS	INF	ERA	
0	1	1	2	13	2013	3.836593	26.861971	0.428889	3.196428	12.076
1	1	1	2	13	2014	5.599992	29.965430	-0.548758	3.045024	8.884
2	1	1	2	13	2015	4.266334	30.984761	0.218786	2.394557	8.583
3	1	1	2	13	2016	4.580100	29.832095	-0.794050	3.712403	5.720
4	1	1	2	13	2017	7.329021	28.630991	1.769412	3.833422	6.256

```
In [25]: 1 # Binning Zscore into categories (example: high risk = 1, low risk = 0)
2 df['Risk_Category'] = pd.cut(df['Zscore'], bins=[-np.inf, 1.8, np.inf], labels=[1, 0])
```

Purpose

The purpose of this line is to transform a continuous measure (Zscore) into a binary categorical variable (Risk_Category), which can be useful for classification tasks. In this case:

- A Zscore of 1.8 or below is labeled as 1 (high risk).
- A Zscore above 1.8 is labeled as 0 (low risk).

```
In [26]: 1 # Features and target
2 features = df.drop(columns=['Zscore', 'id', 'Countries', 'Banks', 'Risk_Category'])
3 target = df['Risk_Category']
```

```
In [27]: 1 # Splitting the data
2 X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)
```

```
In [28]: 1 # Creating the preprocessing
2 numeric_features = features.select_dtypes(include=['int64', 'float64']).columns
3 numeric_transformer = StandardScaler()
4
5 preprocessor = ColumnTransformer(
6     transformers=[
7         ('num', numeric_transformer, numeric_features),
8     ])
```

```
In [29]: 1 # Creating the pipeline
2
3 from sklearn.ensemble import VotingClassifier
4
5 # Creating individual models
6 rf_model = Pipeline(steps=[('preprocessor', preprocessor),
7                             ('classifier', RandomForestClassifier(n_estimators=100)),
8
9 gb_model = Pipeline(steps=[('preprocessor', preprocessor),
10                             ('classifier', GradientBoostingClassifier(n_estimators=100)),
11
12 # Creating the voting classifier
13 voting_clf = VotingClassifier(estimators=[('rf', rf_model), ('gb', gb_model)], voting='hard')
14
```

```
In [30]: 1 # Fitting the model
2 voting_clf.fit(X_train, y_train)
```

```
Out[30]: VotingClassifier(estimators=[('rf',
                                     Pipeline(steps=[('preprocessor',
                                                         ColumnTransformer(transformers=[('num',
                                                                 StandardScaler(),
                                                                 Index(['Countries_Num', 'Year', 'RIR', 'SFS', 'INF', 'ERA', 'INL', 'DEBT',
                                                                 'SIZE', 'CC', 'GE', 'PS', 'RQ', 'RL', 'VA'],
                                                                 dtype='object'))])),
                                                         ('classifier',
                                                         RandomForestClassifier(random_state=42))])),
                                     ('gb',
                                     Pipeline(steps=[('preprocessor',
                                                         ColumnTransformer(transformers=[('num',
                                                                 StandardScaler(),
                                                                 Index(['Countries_Num', 'Year', 'RIR', 'SFS', 'INF', 'ERA', 'INL', 'DEBT',
                                                                 'SIZE', 'CC', 'GE', 'PS', 'RQ', 'RL', 'VA'],
                                                                 dtype='object'))])),
                                                         ('classifier',
                                                         GradientBoostingClassifier(random_state=42))])),
                                     voting='soft'])
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [31]: 1 # Predictions
2 y_pred_train = voting_clf.predict(X_train)
3 y_pred_test = voting_clf.predict(X_test)
```

```
In [32]: 1 # Evaluation
2 print("Training Classification Report:")
3 print(classification_report(y_train, y_pred_train))
4
5 print("*****"*8)
6 print("Test Classification Report:")
7 print(classification_report(y_test, y_pred_test))
8
```

Training Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	374
1	1.00	1.00	1.00	219
accuracy			1.00	593
macro avg	1.00	1.00	1.00	593
weighted avg	1.00	1.00	1.00	593

Test Classification Report:

	precision	recall	f1-score	support
0	0.88	0.91	0.89	87
1	0.86	0.82	0.84	62
accuracy			0.87	149
macro avg	0.87	0.87	0.87	149
weighted avg	0.87	0.87	0.87	149


```
In [33]: 1 # Confusion Matrix
2 print("Confusion Matrix:")
3 cf_matrix = confusion_matrix(y_test, y_pred_test)
4 print(cf_matrix)
```

Confusion Matrix:

```
[[79  8]
 [11 51]]
```

```
In [34]: 1 import seaborn as sns  
2 sns.heatmap(cf_matrix, annot=True)
```

Out[34]: <Axes: >



```
In [ ]: 1
```